

(https://databricks.com)

```
# To import mlflow and autolog machine learning runs
```

```
import mlflow
mlflow.pySpark.ml.autolog()
```

```
# The FaultDataset in DBFS
```

```
dbutils.fs.ls("/FileStore/tables/")
```

```
Out[2]: [FileInfo(path='dbfs:/FileStore/tables/FaultDataset.csv', name='FaultDataset.csv', size=1703184, modificationTime=1678589742000),
  FileInfo(path='dbfs:/FileStore/tables/Occupancy_Detection_Data.csv', name='Occupancy_Detection_Data.csv', size=50968, modificationTime=1677673883000),
  FileInfo(path='dbfs:/FileStore/tables/account-models/', name='account-models/', size=0, modificationTime=0),
  FileInfo(path='dbfs:/FileStore/tables/accounts/', name='accounts/', size=0, modificationTime=0),
  FileInfo(path='dbfs:/FileStore/tables/accounts.zip', name='accounts.zip', size=5297592, modificationTime=1675260543000),
  FileInfo(path='dbfs:/FileStore/tables/activations/', name='activations/', size=0, modificationTime=0),
  FileInfo(path='dbfs:/FileStore/tables/activations.zip', name='activations.zip', size=8411369, modificationTime=1675257654000),
  FileInfo(path='dbfs:/FileStore/tables/clinicaltrial_2019/', name='clinicaltrial_2019/', size=0, modificationTime=0),
  FileInfo(path='dbfs:/FileStore/tables/clinicaltrial_2019.csv', name='clinicaltrial_2019.csv', size=42400056, modificationTime=1679426915000),
  FileInfo(path='dbfs:/FileStore/tables/clinicaltrial_2019.zip', name='clinicaltrial_2019.zip', size=9707871, modificationTime=1678924265000),
  FileInfo(path='dbfs:/FileStore/tables/clinicaltrial_2020/', name='clinicaltrial_2020/', size=0, modificationTime=0),
  FileInfo(path='dbfs:/FileStore/tables/clinicaltrial_2020.csv', name='clinicaltrial_2020.csv', size=46318151, modificationTime=1679431882000),
  FileInfo(path='dbfs:/FileStore/tables/clinicaltrial_2020.zip', name='clinicaltrial_2020.zip', size=10599182, modificationTime=1678915950000),
  FileInfo(path='dbfs:/FileStore/tables/clinicaltrial_2021/', name='clinicaltrial_2021/', size=0, modificationTime=0),
```

```
# To read the data we want to work with into spark DataFrame
```

```
FaultsDF = spark.read.csv("/FileStore/tables/FaultDataset.csv",
                          header = "true",
                          inferSchema = "true")
```

```
# To have an overview of what the data looks like
```

```
FaultsDF.display()
```

Table									
	1	2	3	4	5	6	7	8	9
1	0.3503125	0.3496875	0.35	0.3459375	0.3475	0.3459375	0.341875	0.3434375	0.355
2	0.5090625	0.484375	0.046875	0.071875	0.06	0.0634375	0.0575	0.0546875	0.0559375
3	0.0928125	0.0975	0.1096875	0.1025	0.09625	0.1053125	0.09875	0.098125	0.091875
4	0.09375	0.089375	0.091875	0.0996875	0.0909375	0.096875	0.0940625	0.096875	0.096875
5	0.036875	0.0440625	0.038125	0.0428125	0.0353125	0.0340625	0.033125	0.0403125	0.0346875
6	0.135625	0.3034375	0.13875	0.140625	0.126875	0.130625	0.139375	0.143125	0.1290625
7	0.3446875	0.35125	0.3353125	0.3471875	0.34625	0.348125	0.3478125	0.3521875	0.3525

9,292 rows

```
#To convert to pandas dataframe so that we can perform some EDA and visualisations
```

```
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
```

```
# To convert FaultsDF To pandas dataframe
```

```
FaultPandasFrame = FaultsDF.toPandas()
```


```
# To see the first five rows in the data
```

```
FaultPandasFrame.head()
```

	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
0	0.350313	0.349687	0.350000	0.345938	0.347500	0.345938	0.341875	0.343438	0.355000	0.355312	0.345938	0.352500	0.357500	0.359063	0.358750
1	0.509062	0.484375	0.046875	0.071875	0.060000	0.063437	0.057500	0.054688	0.055938	0.058125	0.062812	0.065625	0.064062	0.063437	0.053437
2	0.092813	0.097500	0.109687	0.102500	0.096250	0.105313	0.098750	0.098125	0.091875	0.090938	0.098750	0.103125	0.100000	0.103438	0.101562
3	0.093750	0.089375	0.091875	0.099687	0.090938	0.096875	0.094062	0.096875	0.096875	0.099375	0.099375	0.095937	0.095937	0.094062	0.091250
4	0.036875	0.044062	0.038125	0.042813	0.035312	0.034063	0.033125	0.040313	0.034688	0.036875	0.035625	0.036250	0.040938	0.039375	0.035000

```
# To see the class balance using a histogram plot.
# The result shows that we have a balance class between fault detected or not
sns.histplot(FaultPandasFrame, x='fault_detected')
```

Out[8]: <AxesSubplot:xlabel='fault_detected', ylabel='Count'>



FaultPandasFrame.shape

Out[9]: (9292, 21)

```
# To get a summary of the data
FaultPandasFrame.describe()
```

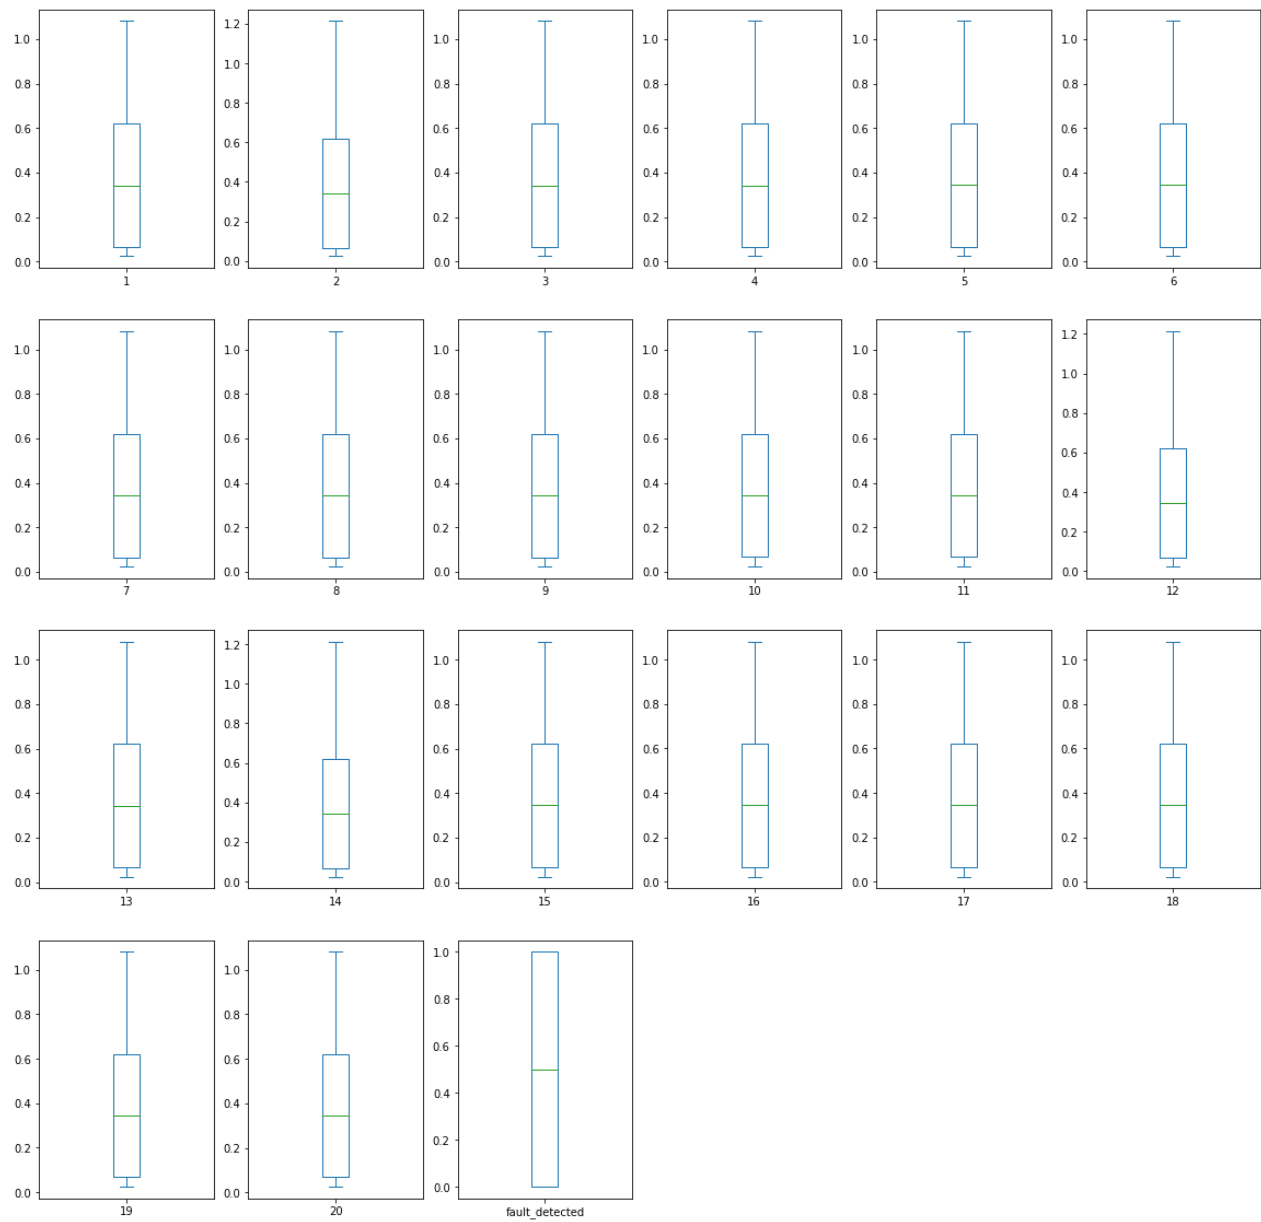
	1	2	3	4	5	6	7	8	9	10	11
count	9292.000000	9292.000000	9292.000000	9292.000000	9292.000000	9292.000000	9292.000000	9292.000000	9292.000000	9292.000000	9292.000000
mean	0.341623	0.342631	0.342121	0.342139	0.342843	0.342828	0.342715	0.343066	0.343173	0.343925	0.344108
std	0.289195	0.289088	0.289164	0.289164	0.288965	0.289089	0.289195	0.289192	0.289340	0.289012	0.289200
min	0.024375	0.024375	0.024375	0.024375	0.024375	0.024375	0.024375	0.024375	0.024375	0.024375	0.024375
25%	0.064062	0.064609	0.064375	0.064375	0.065000	0.065000	0.064687	0.065312	0.065000	0.065937	0.066172
50%	0.342187	0.343750	0.342813	0.342031	0.343594	0.343750	0.342813	0.343125	0.342813	0.345000	0.345313
75%	0.618437	0.619062	0.619062	0.619062	0.619062	0.619062	0.619141	0.619687	0.619766	0.620000	0.620313
max	1.080938	1.213437	1.080938	1.080938	1.080938	1.080938	1.080938	1.080938	1.080938	1.080938	1.080938

```
#Histogram plot of the various features in the data to see how the vibration sensor readings from the data are distributed.
```

```
FaultPandasFrame.hist()
plt.gcf().set_size_inches(20,20)
plt.show()
```



```
# Using a box and whisker plot to check if there are outliers in the virbration readings
FaultPandasFrame.plot(kind='box', subplots=True, layout=(4,6))
plt.gcf().set_size_inches(20,20)
plt.show()
```



FaultPandasFrame.corr()

	1	2	3	4	5	6	7	8	9	10	11	12	13	
1	1.000000	0.843847	0.815780	0.826661	0.819727	0.806820	0.800609	0.800667	0.785560	0.781098	0.781228	0.772658	0.773191	0.778808
2	0.843847	1.000000	0.843153	0.813794	0.824917	0.818478	0.806115	0.799769	0.798175	0.785856	0.778867	0.780272	0.771881	0.772710
3	0.815780	0.843153	1.000000	0.844225	0.815464	0.826045	0.818762	0.806707	0.799662	0.801067	0.785955	0.780895	0.780712	0.774210
4	0.826661	0.813794	0.844225	1.000000	0.844152	0.815404	0.825898	0.821163	0.807515	0.801069	0.799968	0.786714	0.783222	0.783210
5	0.819727	0.824917	0.815464	0.844152	1.000000	0.843323	0.814768	0.827324	0.820063	0.807351	0.800829	0.801012	0.786438	0.783210
6	0.806820	0.818478	0.826045	0.815404	0.843323	1.000000	0.842480	0.813768	0.824977	0.821148	0.807483	0.799901	0.800890	0.787610
7	0.800609	0.806115	0.818762	0.825898	0.814768	0.842480	1.000000	0.844601	0.814972	0.826524	0.819826	0.807104	0.801846	0.802110
8	0.800667	0.799769	0.806707	0.821163	0.827324	0.813768	0.844601	1.000000	0.843669	0.813676	0.824552	0.819180	0.807114	0.801910
9	0.785560	0.798175	0.799662	0.807515	0.820063	0.824977	0.814972	0.843669	1.000000	0.843174	0.813524	0.825630	0.820179	0.807510
10	0.781098	0.785856	0.801067	0.801069	0.807351	0.821148	0.826524	0.813676	0.843174	1.000000	0.842913	0.812985	0.825341	0.819510
11	0.781228	0.778867	0.785955	0.799968	0.800829	0.807483	0.819826	0.824552	0.813524	0.842913	1.000000	0.842194	0.813314	0.824610
12	0.772658	0.780272	0.780895	0.786714	0.801012	0.799901	0.807104	0.819180	0.825630	0.812985	0.842194	1.000000	0.843725	0.813510
13	0.773191	0.771881	0.780712	0.783222	0.786438	0.800890	0.801846	0.807114	0.820179	0.825341	0.813314	0.843725	1.000000	0.842610
14	0.778808	0.772782	0.774265	0.783219	0.783206	0.787682	0.802103	0.801997	0.807529	0.819571	0.824605	0.813560	0.842642	1.000000
15	0.770734	0.777026	0.772514	0.774887	0.782338	0.782074	0.788105	0.800667	0.801262	0.807342	0.818214	0.825275	0.813118	0.842510
16	0.758679	0.769570	0.777459	0.774451	0.774300	0.781673	0.782957	0.788960	0.801438	0.801288	0.806279	0.819909	0.826241	0.812510
17	0.765337	0.757522	0.770544	0.780225	0.775226	0.775094	0.783322	0.782572	0.789297	0.801063	0.800629	0.807568	0.821294	0.825610
18	0.757268	0.764096	0.759063	0.771988	0.779908	0.774694	0.775707	0.783127	0.782368	0.788638	0.800450	0.801630	0.807897	0.820210
19	0.752535	0.758217	0.766772	0.761376	0.772765	0.780398	0.776955	0.777520	0.784443	0.782847	0.789137	0.802929	0.802816	0.807910
20	0.751781	0.751194	0.757990	0.766829	0.760875	0.771479	0.780377	0.776530	0.776484	0.783898	0.781167	0.788763	0.801218	0.802010
fault_detected	0.667028	0.666130	0.668463	0.670146	0.669696	0.669190	0.671794	0.673588	0.674972	0.675330	0.675667	0.677753	0.681194	0.681710


Correlation matrix of the readings from the vibration sensors and the fault detection

```
data = FaultPandasFrame
```

```
plt.figure(figsize=(20,8))
```

```
sns.heatmap(data.corr(), annot = True)
```

Out[14]: <AxesSubplot:>




#Scatter plot matrix to show the relationship between two variables of the parameters being read by the vibration sensor

```
from pandas.plotting import scatter_matrix
```

```
scatter_matrix(FaultPandasFrame)
```

```
plt.gcf().set_size_inches(20,20)
```

```
plt.show()
```



```
df2 = FaultPandasFrame
```

```
#import pandas as pd
```

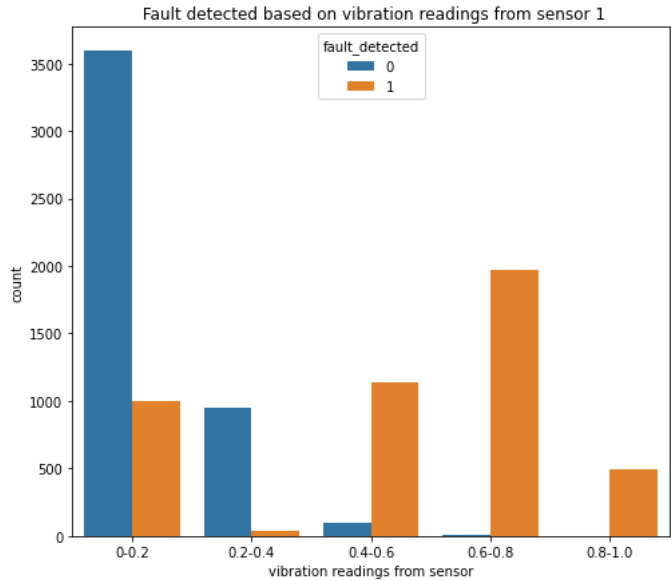
```
# We want to visualise the fault detection based on the vibration readings from each features
# To get a better visualisation, the vibration readings will be put into bins such that we capture from the minimum
reading and the maximum reading
# First we can introduce a variable to hold the pandas dataframe create earlier
df = FaultPandasFrame
```

```
df.describe()
```


	1	2	3	4	5	6	7	8	9	10	11
count	9292.000000	9292.000000	9292.000000	9292.000000	9292.000000	9292.000000	9292.000000	9292.000000	9292.000000	9292.000000	9292.000000
mean	0.341623	0.342631	0.342121	0.342139	0.342843	0.342828	0.342715	0.343066	0.343173	0.343925	0.344108
std	0.289195	0.289088	0.289164	0.289164	0.288965	0.289089	0.289195	0.289192	0.289340	0.289012	0.289200
min	0.024375	0.024375	0.024375	0.024375	0.024375	0.024375	0.024375	0.024375	0.024375	0.024375	0.024375
25%	0.064062	0.064609	0.064375	0.064375	0.065000	0.065000	0.064687	0.065312	0.065000	0.065937	0.066172
50%	0.342187	0.343750	0.342813	0.342031	0.343594	0.343750	0.342813	0.343125	0.342813	0.345000	0.345313
75%	0.618437	0.619062	0.619062	0.619062	0.619062	0.619062	0.619141	0.619687	0.619766	0.620000	0.620313
max	1.080938	1.213437	1.080938	1.080938	1.080938	1.080938	1.080938	1.080938	1.080938	1.080938	1.080938

```
# Putting sensor readings into bins
bin_sensor_1 = [0, 0.2, 0.4, 0.6, 0.8, 1.0]
labels = ['0-0.2', '0.2-0.4', '0.4-0.6', '0.6-0.8', '0.8-1.0']
df['binned_1'] = pd.cut(df['1'], bins=bin_sensor_1 , labels =labels)
plt.figure(figsize=(8,7))
sns.countplot(x='binned_1',hue='fault_detected', data = df)
plt.title(" Fault detected based on vibration readings from sensor 1", fontsize = 12)
plt.xlabel('vibration readings from sensor')
```

```
Out[20]: Text(0.5, 0, 'vibration readings from sensor')
```



```
plt.figure(figsize=(18,6))
sns.set(font_scale=1.2)
ax=sns.countplot(x='binned_1',hue='fault_detected', data = df)
ax.set_title('Fault detected based on vibration readings from sensor 1' , fontsize = 20)
plt.xlabel('vibration readings from sensor', fontsize=15)
plt.ylabel('Count ', fontsize=15)
plt.xticks(rotation='horizontal')
for p in ax.patches:
    for p in ax.patches:
        ax.annotate(format(p.get_height(), '.0f'),
                    (p.get_x()+0.3, p.get_height()), ha='center', va='bottom', color='black', size=15)
```




```
# Cut vibration readings into bin and show the percentage of fault detected within each bin
bin_sen_1 = pd.cut(df['1'], bins=5, ordered=True)
sen1_labels = [f'({interval.left}, {interval.right})' for interval in bin_sen_1.cat.categories]
df['bins_sen_1'] = bin_sen_1.cat.rename_categories(sen1_labels)

# Create countplot
plt.figure(figsize=(18, 6))
sns.set(font_scale=1.2)
ax = sns.countplot(x='bins_sen_1', hue='fault_detected', data= df)
ax.set_title('Percentage of fault detected by vibrations from sensor 1', fontsize=15)
plt.xlabel('Vibration', fontsize=12)
plt.ylabel('count', fontsize=12)
plt.xticks(rotation='horizontal')

# Annotate plot with percentages
total_counts = df.groupby('bins_sen_1')['fault_detected'].count().reset_index(name='count')
fault_counts = df.groupby(['bins_sen_1', 'fault_detected']).size().reset_index(name='fault_count')

for p, label in zip(ax.patches, ax.get_xticklabels()):
    height = p.get_height()
    bin_label = label.get_text()
    bin_counts = fault_counts[fault_counts['bins_sen_1'] == bin_label]['fault_count']
    total_count = total_counts[total_counts['bins_sen_1'] == bin_label]['count'].iloc[0]
    percentages = [count / total_count * 100 for count in bin_counts]
    percentage_str = ', '.join([f'{percentage:.2f}%' for percentage in percentages])
    ax.annotate(percentage_str, (p.get_x() + p.get_width() / 2, height),
                ha='center', va='bottom', color='black', size=12,
                xytext=(28, 0), textcoords='offset points')

plt.show()
```





```
# The percentage of fault within each bins
bins = [0, 0.2, 0.4, 0.6, 0.8, 1.0]
labels = ['0-0.2', '0.2-0.4', '0.4-0.6', '0.6-0.8', '0.8-1.0']
df['binned'] = pd.cut(df['1'], bins=bins, labels =labels)

# group by binned and calculate percentage of each fault_detected
grouped = df.groupby('binned')['fault_detected'].value_counts(normalize=True).reset_index(name='percentage')
plt.figure(figsize=(18, 6))

sns.barplot(x='binned',y='percentage', hue='fault_detected', data = grouped)
plt.title("The percentage of fault detected within each bins in sensor 1", fontsize = 12)
plt.xlabel('vibration readings from sensor 1')

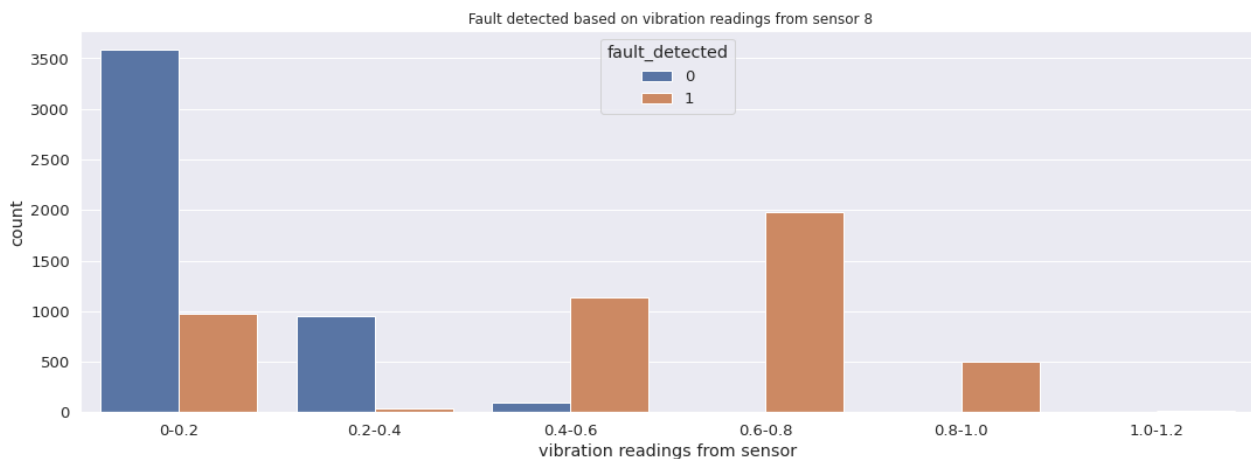
Out[23]: Text(0.5, 0, 'vibration readings from sensor 1')
```




```
#Let's visualise for some other sensors. We can choose sensor 8
bins = [0, 0.2, 0.4, 0.6, 0.8, 1.0, 1.2]
labels = ['0-0.2', '0.2-0.4', '0.4-0.6', '0.6-0.8', '0.8-1.0', '1.0-1.2']

df['binned'] = pd.cut(df['8'], bins=bins, labels =labels)
plt.figure(figsize=(18,6))
sns.countplot(x='binned',hue='fault_detected', data = df)
plt.title(" Fault detected based on vibration readings from sensor 8", fontsize = 12)
plt.xlabel('vibration readings from sensor 8')

Out[24]: Text(0.5, 0, 'vibration readings from sensor 8')
```



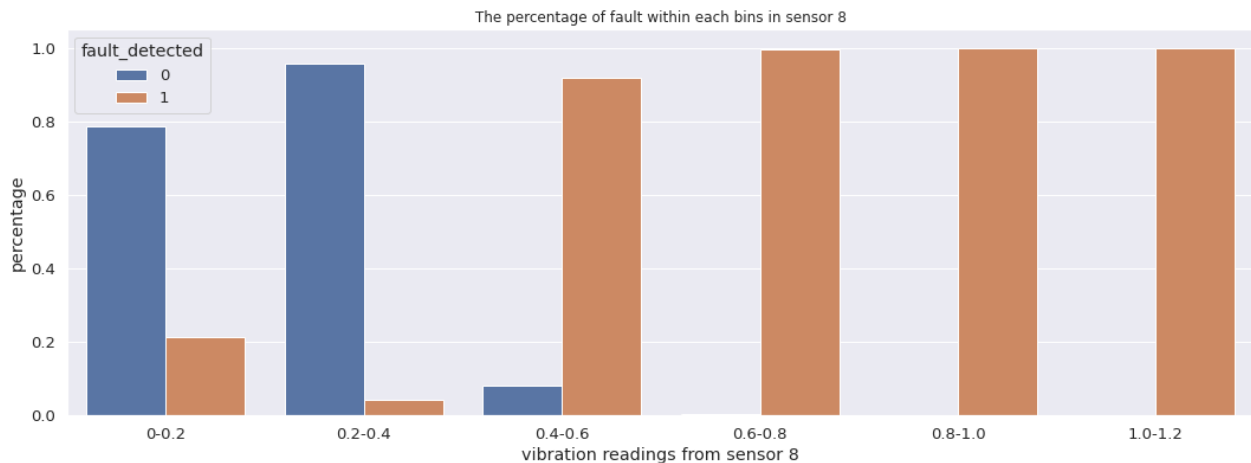
```
plt.figure(figsize=(18,6))
sns.set(font_scale=1.2)
ax=sns.countplot(x='binned',hue='fault_detected', data = df)
ax.set_title('Fault detected based on vibration readings from sensor 8' , fontsize = 20)
plt.xlabel('vibration readings from sensor', fontsize=15)
plt.ylabel('Count ', fontsize=15)
plt.xticks(rotation='horizontal')
for p in ax.patches:
    for p in ax.patches:
        ax.annotate(format(p.get_height(), '.0f'),
                    (p.get_x()+0.3, p.get_height()), ha='center', va='bottom', color='black', size=15)
```



```
# The percentage of fault within each bins
#bins = [0, 0.25, 0.5, 0.75, 1.0, 1.25]
#labels = ['0-0.25', '0.25-0.50', '0.50-0.75', '0.75-1.0', '1.0-1.25']
bins = [0, 0.2, 0.4, 0.6, 0.8, 1.0, 1.2]
labels = ['0-0.2', '0.2-0.4', '0.4-0.6', '0.6-0.8', '0.8-1.0', '1.0-1.2']
df['binned'] = pd.cut(df['8'], bins=bins, labels =labels)
# group by binned and calculate percentage of each fault_detected
grouped = df.groupby('binned')['fault_detected'].value_counts(normalize=True).reset_index(name='percentage')
```

```
plt.figure(figsize=(18,6))
sns.barplot(x='binned',y='percentage', hue='fault_detected', data = grouped)
plt.title(" The percentage of fault within each bins in sensor 8 ", fontsize = 12)
plt.xlabel('vibration readings from sensor 8')
```

Out[26]: Text(0.5, 0, 'vibration readings from sensor 8')



```

# Cut vibration readings into bin and show the percentage of fault detected within each bin
bin_sen_8 = pd.cut(df['8'], bins=5, ordered=True)
sen8_labels = [f'({interval.left}, {interval.right})' for interval in bin_sen_8.cat.categories]
df['bins_sen_8'] = bin_sen_8.cat.rename_categories(sen8_labels)


# Create countplot
plt.figure(figsize=(18, 6))
sns.set(font_scale=1.2)
ax = sns.countplot(x='bins_sen_8', hue='fault_detected', data= df)
ax.set_title('Percentage of fault detected by vibrations from sensor 8', fontsize=15)
plt.xlabel('Vibration', fontsize=12)
plt.ylabel('count', fontsize=12)
plt.xticks(rotation='horizontal')

# Annotate plot with percentages
total_counts = df.groupby('bins_sen_8')['fault_detected'].count().reset_index(name='count')
fault_counts = df.groupby(['bins_sen_8', 'fault_detected']).size().reset_index(name='fault_count')

for p, label in zip(ax.patches, ax.get_xticklabels()):
    height = p.get_height()
    bin_label = label.get_text()
    bin_counts = fault_counts[fault_counts['bins_sen_8'] == bin_label]['fault_count']
    total_count = total_counts[total_counts['bins_sen_8'] == bin_label]['count'].iloc[0]
    percentages = [count / total_count * 100 for count in bin_counts]
    percentage_str = ', '.join([f'{percentage:.2f}%' for percentage in percentages])
    ax.annotate(percentage_str, (p.get_x() + p.get_width() / 2, height),
                ha='center', va='bottom', color='black', size=12,
                xytext=(28, 0), textcoords='offset points')

plt.show()

```

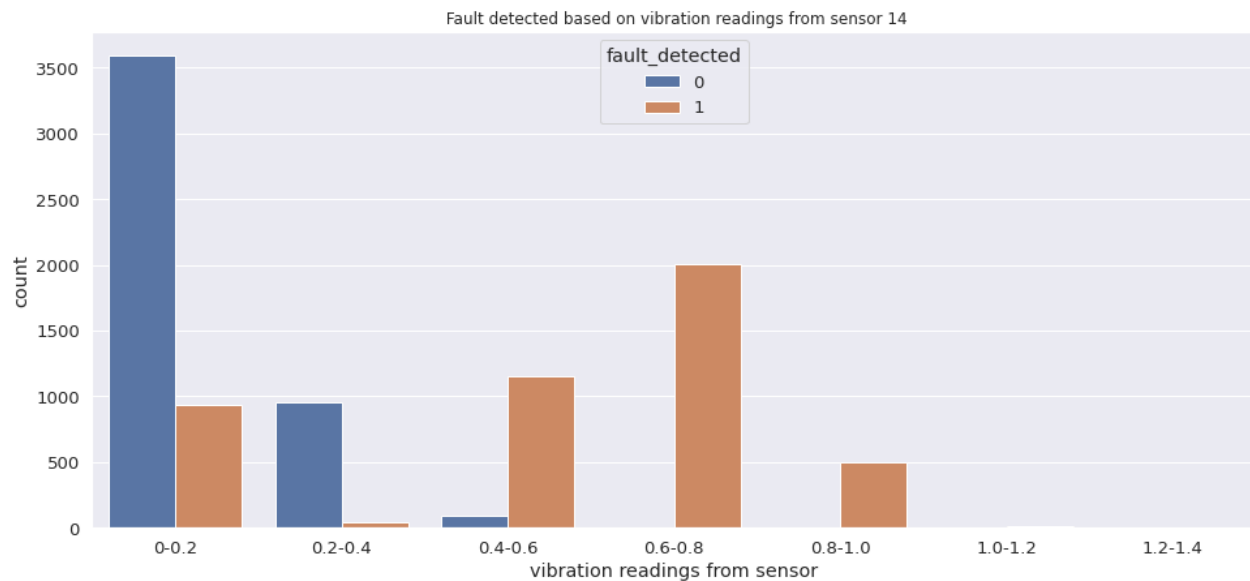


```


bins = [0, 0.2, 0.4, 0.6, 0.8, 1.0, 1.2, 1.4]
labels = ['0-0.2', '0.2-0.4', '0.4-0.6', '0.6-0.8', '0.8-1.0', '1.0-1.2', '1.2-1.4']
df['binned_14'] = pd.cut(df['14'], bins=bins, labels =labels)
plt.figure(figsize=(16,7))
sns.countplot(x='binned_14', hue='fault_detected', data= df)
plt.title(" Fault detected based on vibration readings from sensor 14", fontsize = 12)
plt.xlabel('vibration readings from sensor')

```

Out[28]: Text(0.5, 0, 'vibration readings from sensor')



```
plt.figure(figsize=(18,6))
sns.set(font_scale=1.2)
ax=sns.countplot(x='binned_14',hue='fault_detected', data = df)
ax.set_title('Fault detected based on vibration readings from sensor 14' , fontsize = 20)
plt.xlabel('vibration readings from sensor', fontsize=15)
plt.ylabel('Count ', fontsize=15)
plt.xticks(rotation='horizontal')
for p in ax.patches:
    for p in ax.patches:
        ax.annotate(format(p.get_height(), '.0f'),
                    (p.get_x()+0.3, p.get_height()), ha='center', va='bottom', color='black', size=15)
```



```

# Cut vibration readings into bin and show the percentage of fault detected within each bin
bin_sen_14 = pd.cut(df['14'], bins=6, ordered=True)
sen14_labels = [f'({interval.left}, {interval.right})' for interval in bin_sen_14.cat.categories]
df['bins_sen_14'] = bin_sen_14.cat.rename_categories(sen14_labels)


# Create countplot
plt.figure(figsize=(18, 6))
sns.set(font_scale=1.2)
ax = sns.countplot(x='bins_sen_14', hue='fault_detected', data= df)
ax.set_title('Percentage of fault detected by vibrations from sensor 14', fontsize=15)
plt.xlabel('Vibration', fontsize=12)
plt.ylabel('count', fontsize=12)
plt.xticks(rotation='horizontal')

# Annotate plot with percentages
total_counts = df.groupby('bins_sen_14')['fault_detected'].count().reset_index(name='count')
fault_counts = df.groupby(['bins_sen_14', 'fault_detected']).size().reset_index(name='fault_count')

for p, label in zip(ax.patches, ax.get_xticklabels()):
    height = p.get_height()
    bin_label = label.get_text()
    bin_counts = fault_counts[fault_counts['bins_sen_14'] == bin_label]['fault_count']
    total_count = total_counts[total_counts['bins_sen_14'] == bin_label]['count'].iloc[0]
    percentages = [count / total_count * 100 for count in bin_counts]
    percentage_str = ', '.join([f'{percentage:.2f}%' for percentage in percentages])
    ax.annotate(percentage_str, (p.get_x() + p.get_width() / 2, height),
               ha='center', va='bottom', color='black', size=12,
               xytext=(28, 0), textcoords='offset points')

plt.show()

```

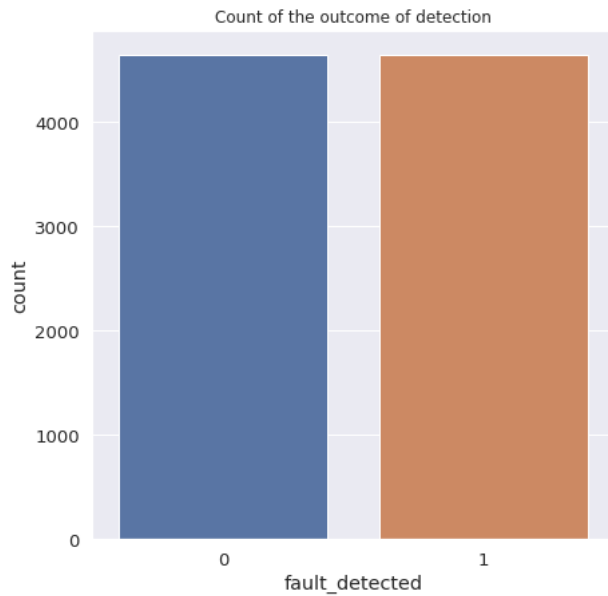


```

plt.figure(figsize=(7,7))
sns.countplot(x='fault_detected', data = df2)
plt.title('Count of the outcome of detection', fontsize = 12)

```

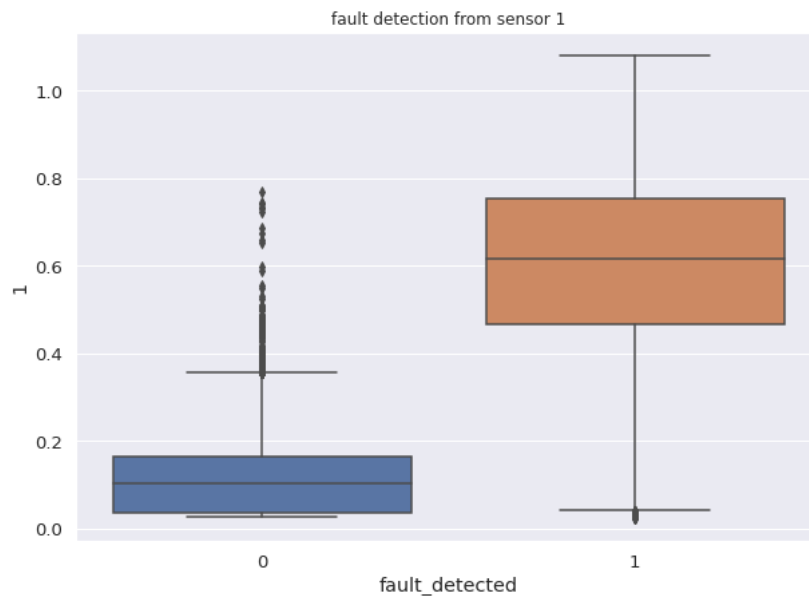
Out[31]: Text(0.5, 1.0, 'Count of the outcome of detection')



```
#To check if their are duplicated values
#set(FaultPandasFrame.duplicated())
```


```
# A box plot to show the outcome of fault detection based on the readings for each sensor
# Readings from few sensors were checked: sensor 1, sensor 8 and sensor 14
plt.figure(figsize=(10,7))
sns.boxplot(y='1', x='fault_detected', data=df2)
plt.title('fault detection from sensor 1', fontsize = 12)
```

```
Out[33]: Text(0.5, 1.0, 'fault detection from sensor 1')
```



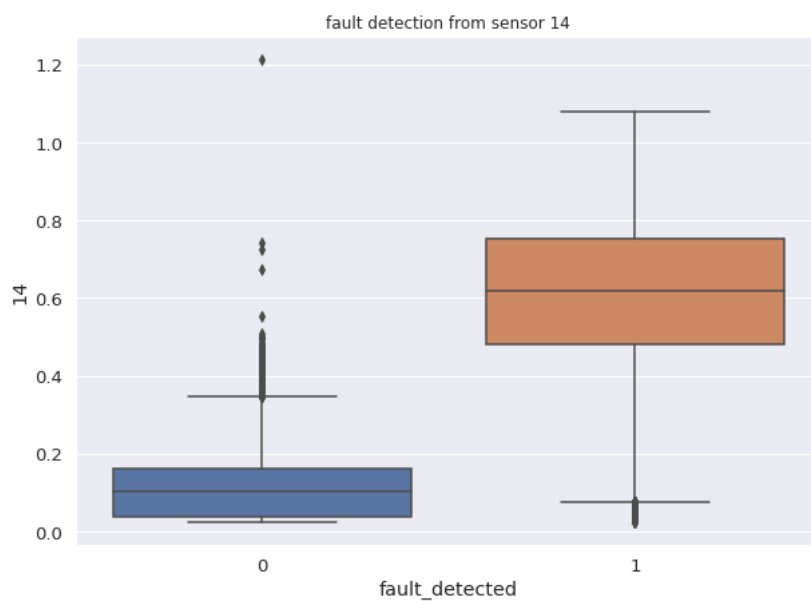
```
plt.figure(figsize=(10,7))
sns.boxplot(y='8', x='fault_detected', data=df2)
plt.title('fault detection from sensor 8', fontsize = 12)
```

```
Out[34]: Text(0.5, 1.0, 'fault detection from sensor 8')
```



```
plt.figure(figsize=(10,7))
sns.boxplot(y='14', x='fault_detected', data=df2)
plt.title('fault detection from sensor 14', fontsize = 12)
```

Out[35]: Text(0.5, 1.0, 'fault detection from sensor 14')



```
# Using the DataFrame created initially for further EDA
FaultsDF.printSchema()
```

```
root
|-- 1: double (nullable = true)
|-- 2: double (nullable = true)
|-- 3: double (nullable = true)
|-- 4: double (nullable = true)
|-- 5: double (nullable = true)
|-- 6: double (nullable = true)
|-- 7: double (nullable = true)
|-- 8: double (nullable = true)
|-- 9: double (nullable = true)
|-- 10: double (nullable = true)
|-- 11: double (nullable = true)
```

```
|-- 12: double (nullable = true)
|-- 13: double (nullable = true)
|-- 14: double (nullable = true)
|-- 15: double (nullable = true)
|-- 16: double (nullable = true)
|-- 17: double (nullable = true)
|-- 18: double (nullable = true)
|-- 19: double (nullable = true)
|-- 20: double (nullable = true)
```

```
FaultsDF.describe().display()
```

Table						
	summary ▲	1 ▲	2 ▲	3 ▲	4 ▲	5
1	count	9292	9292	9292	9292	9292
2	mean	0.34162330499354226	0.34263116121394677	0.3421213812957383	0.34213907124407966	0.342843440594058
3	stddev	0.28919489486260785	0.2890875372793958	0.28916422490616933	0.28916356333107296	0.288964655440387
4	min	0.024375	0.024375	0.024375	0.024375	0.024375
5	max	1.0809375	1.2134375	1.0809375	1.0809375	1.0809375

```
# Let's also see the distribution of the result from the fault detected if the class is balance or not
# here the count function will be imported so that we can use it to count to outcome of each class of result
from pyspark.sql.functions import count as _count
```

```
# To check Class balance
FaultsDF.groupBy("fault_detected").agg(_count("*")).display()
```

Table			
	fault_detected ▲	count(1) ▲	
1	1	4646	
2	0	4646	
2 rows			

```
# Importing some functions to help with further EDA
from pyspark.sql.functions import isnan, when, count, col
```

```
# Checking if there are missing values in each column
FaultsDF.select([_count(when(isnan(c), c)).alias(c) for c in FaultsDF.columns]).display()
```

Table										
	1	2	3	4	5	6	7	8	9	
1	0	0	0	0	0	0	0	0	0	
1 row										

```
# Checking if there are missing values in each column
FaultsDF.select([_count(when(isnan(c), c)).alias(c) for c in FaultsDF.columns]).show()
```

[illegible]


```
duplicates_ = FaultsDF.groupBy(FaultsDF.columns).count().filter(col("count") > 1)
print("Duplicates based on all columns:")
duplicates_.display()
```

Duplicates based on all columns:

Table									
	1 ▲	2 ▲	3 ▲	4 ▲	5 ▲	6 ▲	7 ▲	8 ▲	9 ▲
1	0.035625	0.0359375	0.0315625	0.035625	0.0365625	0.03625	0.040625	0.0409375	0.0421875
2	0.09875	0.09875	0.0978125	0.1028125	0.1003125	0.09625	0.1015625	0.10125	0.096875
3	0.0346875	0.036875	0.0353125	0.03875	0.0334375	0.0359375	0.0390625	0.0334375	0.03625
4	0.10625	0.10625	0.1021875	0.0971875	0.101875	0.1065625	0.1040625	0.101875	0.104375
5	0.1009375	0.1040625	0.101875	0.1078125	0.10875	0.1059375	0.1096875	0.111875	0.110625
6	0.035625	0.0334375	0.034375	0.03875	0.03875	0.0359375	0.03375	0.0346875	0.0375
7	0.4028125	0.3915625	0.3915625	0.384375	0.3859375	0.3871875	0.3975	0.398125	0.3896875
300 rows									

```
# To confirm the number of rows and the distinct number of rows so we can identify if there is duplicated rows and see how many they are
row_count = FaultsDF.count()
distinct_count = FaultsDF.distinct().count()
print(row_count)
print(distinct_count)
```

9292
8968

```
print(row_count - distinct_count)
```

324

```
# To remove duplicated rows and keep only distinct rows
```

```
FaultsDF_drop = FaultsDF.distinct()
```

```
# To confirm that we now have 8968 rows since we have dropped the duplicated rows
FaultsDF_drop_count = FaultsDF_drop.distinct().count()
FaultsDF_drop_count
```

Out[47]: 8968

```
FaultsDF_drop.display()
```

Table									
	1 ▲	2 ▲	3 ▲	4 ▲	5 ▲	6 ▲	7 ▲	8 ▲	9 ▲
1	0.3634375	0.3665625	0.370625	0.369375	0.368125	0.3678125	0.3803125	0.37625	0.3809375
2	0.03875	0.03	0.0321875	0.0396875	0.0384375	0.031875	0.0384375	0.0384375	0.034375
3	0.1428125	0.1459375	0.1346875	0.1315625	0.130625	0.1346875	0.155	0.12625	0.1765625
4	0.033125	0.025	0.0303125	0.026875	0.033125	0.0328125	0.025	0.0290625	0.0328125
5	0.03375	0.0390625	0.03125	0.0390625	0.0384375	0.0390625	0.0325	0.03875	0.03625
6	0.0328125	0.0340625	0.031875	0.03625	0.0378125	0.038125	0.0359375	0.035625	0.034375
7	0.033125	0.0378125	0.0340625	0.03875	0.0325	0.0325	0.035	0.03375	0.03125
8,968 rows									

```
FaultsDF_drop.describe().display()
```

Table						
	summary ▲	1 ▲	2 ▲	3 ▲	4 ▲	5
1	count	8968	8968	8968	8968	8968
2	mean	0.3487404870093662	0.3496980932203395	0.34913992807760935	0.3491733455062442	0.3500254028211428
3	stddev	0.29091106336065264	0.2907930998830016	0.29090468940487874	0.29089397507061654	0.2906356318628472
4	min	0.024375	0.024375	0.024375	0.024375	0.024375
5	max	1.0809375	1.2134375	1.0809375	1.0809375	1.0809375
5 rows						

```

from pyspark.ml.feature import RFormula
Process = RFormula(formula= "fault_detected ~ .")
FaultDF = Process.fit(FaultsDF_drop).transform(FaultsDF_drop)
FaultDF.show(10)

```

2023/05/29 12:47:09 INFO mlflow.utils.autologging_utils: Created MLflow autologging run with ID '59347b1de50248eab25fecf85513b4a4', which will track hyperparameters, performance metrics, model artifacts, and lineage information for the current pyspark.ml workflow

```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+
|      1|      2|      3|      4|      5|      6|      7|      8|      9|     10|     11|     12|
13|     14|     15|     16|     17|     18|     19|    20|fault_detected|          features|label|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+
|0.3634375|0.3665625| 0.370625| 0.369375| 0.368125|0.3678125|0.3803125| 0.37625|0.3809375|0.3815625|0.3928125|0.3834375|
|0.36625| 0.379375|0.3684375|0.3728125| 0.36625| 0.379375| 0.37375| 0.376875|          0|[0.3634375,0.3665...| 0.0|
| 0.03875|      0.03|0.0321875|0.0396875|0.0384375| 0.031875|0.0384375|0.0384375| 0.034375|      0.035|0.0371875|0.0315625|
|0.0328125| 0.031875|      0.0325| 0.03875|0.0365625|0.0378125|      0.035|0.0353125|          0|[0.03875,0.03,0.0...| 0.0
|
|0.1428125|0.1459375|0.1346875|0.1315625| 0.130625|0.1346875|      0.155| 0.12625|0.1765625|0.1403125|0.1434375| 0.139375|
|0.139375| 0.13375| 0.1425| 0.1375| 0.12875| 0.131875| 0.130625|0.1353125|          0|[0.1428125,0.1459...| 0.0|
| 0.033125|      0.025|0.0303125| 0.026875| 0.033125|0.0328125|      0.025|0.0290625|0.0328125|0.0340625|      0.03| 0.0325|
|0.0346875| 0.03125| 0.02875| 0.031875|0.0259375| 0.033125|      0.03| 0.026875|          0|[0.033125,0.025,0.0...| 0.0
|

```

```
#FaultDF.display(10)
```

```
# To split the data into training and testing data
```

```
(trainingDF, testingDF) = FaultDF.randomSplit([0.75, 0.25], seed =50)
```

```

# Training the model
# Using DecisionTreeClassifier
from pyspark.ml.classification import DecisionTreeClassifier
dt = DecisionTreeClassifier(labelCol="label", featuresCol="features")
model = dt.fit(trainingDF)

```

2023/05/29 12:47:11 INFO mlflow.utils.autologging_utils: Created MLflow autologging run with ID '4e74d60dd047452c9a348bbadf204f86', which will track hyperparameters, performance metrics, model artifacts, and lineage information for the current pyspark.ml workflow

2023/05/29 12:48:25 WARNING mlflow.pyspark.ml: Model inputs contain unsupported Spark data types: [StructField('features', VectorUDT(), True)]. Model signature is not logged.

2023/05/29 12:48:29 INFO mlflow.spark: Inferring pip requirements by reloading the logged model from the databricks artifact repository, which can be time-consuming. To speed up, explicitly specify the conda_env or pip_requirements when calling log_model().

2023/05/29 12:49:28 WARNING mlflow.utils.autologging_utils: MLflow autologging encountered a warning: "/databricks/python/n/lib/python3.9/site-packages/_distutils_hack/__init__.py:30: UserWarning: Setuptools is replacing distutils."

```
# To make predictions with the trained model
```

```

predictions = model.transform(testingDF)
predictions.show()

```

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|      1|      2|      3|      4|      5|      6|      7|      8|      9|     10|     11|     12|
13|     14|     15|     16|     17|     18|     19|    20|fault_detected|          features|label| rawPr
ediction|          probability|prediction|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|  0.02625| 0.025625|0.0290625|0.0265625|0.0353125|0.0259375|0.0328125|0.0340625| 0.03375|0.0309375|0.0309375| 0.034375|
0.030625| 0.029375| 0.029375| 0.026875| 0.030625|0.0284375| 0.03125|0.0271875|          0|[0.02625,0.025625...| 0.0|
[3151.0,149.0]|[0.95484848484848...|          0.0|
|0.0265625|0.0284375| 0.02625| 0.0325|0.0296875|0.0353125|0.0328125|0.0265625| 0.03125|0.0328125| 0.029375|0.0284375|
0.0334375|0.0321875|0.0296875| 0.0325|0.0284375|0.0259375| 0.03125| 0.02875|          1|[0.0265625,0.0284...| 1.0
|[3151.0,149.0]|[0.95484848484848...|          0.0|
| 0.026875|0.0346875| 0.03125|0.0303125|0.0315625| 0.02875|0.0278125| 0.034375| 0.0325| 0.0375| 0.030625| 0.03625|
0.0334375|0.0346875|0.0284375|0.0290625|0.0296875| 0.031875|0.0315625|0.0315625|          0|[0.026875,0.03468...| 0.0
|[3151.0,149.0]|[0.95484848484848...|          0.0|
|0.0278125| 0.0275| 0.02875| 0.035| 0.029375|0.0284375| 0.028125|0.0340625| 0.030625|0.0315625|0.0265625|0.0253125|
0.0259375|0.0284375|0.0296875|0.0340625|0.0328125|0.0328125| 0.0325| 0.025|          1|[0.0278125,0.0275...| 1.0
|[3151.0,149.0]|[0.95484848484848...|          0.0|
```

predictions.show(2)

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|      1|      2|      3|      4|      5|      6|      7|      8|      9|     10|     11|     12|
13|     14|     15|     16|     17|     18|     19|    20|fault_detected|          features|label| rawPredi
ction|          probability|prediction|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|  0.02625| 0.025625|0.0290625|0.0265625|0.0353125|0.0259375|0.0328125|0.0340625|0.03375|0.0309375|0.0309375| 0.034375|
0.030625| 0.029375| 0.029375|0.026875| 0.030625|0.0284375|0.03125|0.0271875|          0|[0.02625,0.025625...| 0.0|[31
51.0,149.0]|[0.95484848484848...|          0.0|
|0.0265625|0.0284375| 0.02625| 0.0325|0.0296875|0.0353125|0.0328125|0.0265625|0.03125|0.0328125| 0.029375|0.0284375|0.
0334375|0.0321875|0.0296875| 0.0325|0.0284375|0.0259375|0.03125| 0.02875|          1|[0.0265625,0.0284...| 1.0|[315
1.0,149.0]|[0.95484848484848...|          0.0|
```

only showing top 2 rows

predictions.display()

Table									
	1	2	3	4	5	6	7	8	9
1	0.02625	0.025625	0.0290625	0.0265625	0.0353125	0.0259375	0.0328125	0.0340625	0.03375
2	0.0265625	0.0284375	0.02625	0.0325	0.0296875	0.0353125	0.0328125	0.0265625	0.03125
3	0.026875	0.0346875	0.03125	0.0303125	0.0315625	0.02875	0.0278125	0.034375	0.0325
4	0.0278125	0.0275	0.02875	0.035	0.029375	0.0284375	0.028125	0.0340625	0.030625
5	0.0278125	0.0284375	0.0271875	0.0284375	0.0309375	0.0309375	0.02875	0.0340625	0.0296875
6	0.0278125	0.03	0.0284375	0.029375	0.02875	0.03125	0.0325	0.0290625	0.02625
7	0.0278125	0.038125	0.036875	0.1553125	0.1696875	0.1809375	0.189375	0.176875	0.1646875

2,255 rows

```
# To evaluate the accuracy of predictions made by the model, we can use the multiclass classification evaluator
from pyspark.ml.evaluation import MulticlassClassificationEvaluator
evaluator = MulticlassClassificationEvaluator(labelCol="label", predictionCol="prediction", metricName="accuracy")
accuracy = evaluator.evaluate(predictions)
print("Accuracy score: %g " % (accuracy))
```

Accuracy score: 0.968071

```
# Using MLflow to track experiment
```

```
from pyspark.ml.tuning import ParamGridBuilder, TrainValidationSplit
from pyspark.ml.evaluation import MulticlassClassificationEvaluator
from pyspark.ml.classification import DecisionTreeClassifier
from pyspark.sql.functions import col
```

```
# start an MLflow run and giving it a name for easy identification in the logs
with mlflow.start_run(run_name="decision_tree_classification") as run:
```

```
    # create an instance of decision tree classifier
    dt = DecisionTreeClassifier(featuresCol="features", labelCol="label")

    # creating a parameter grid of different values
    paramGrid = (ParamGridBuilder()
                 .addGrid(dt.maxDepth, [7,9 ])
                 .addGrid(dt.maxBins, [12,32])
                 .addGrid(dt.impurity, ['gini', 'entropy'])
                 .build())

    # creating a TrainValidationSplit object
    tvs = TrainValidationSplit(estimator=dt,
                              estimatorParamMaps=paramGrid,
                              evaluator=MulticlassClassificationEvaluator()
                              )

    # fit the model
    model = tvs.fit(trainingDF)

    # retrieving a dictionary of the hyperparameters and values of the best model
    best_parameters = model.bestModel.extractParamMap()
    # log the best model parameters
    for param, value in best_parameters.items():

        with mlflow.start_run(nested=True):
            mlflow.log_param("depth", 7)
        with mlflow.start_run(nested=True):
            mlflow.log_param("depth", 9)

    # The best trained model
    best_model_1=model.bestModel

    # using the best model on the testing data and evaluating the accuracy

    evaluator = MulticlassClassificationEvaluator(labelCol="label",
                                                  predictionCol="prediction",
                                                  metricName="accuracy")
    accuracy_1 = evaluator.evaluate(best_model_1.transform(testingDF))

    # log the accuracy of the best model
    mlflow.log_metric("accuracy", accuracy_1)

    # end the MLflow run
    mlflow.end_run()
```

```
2023/05/29 13:10:07 INFO mlflow.spark: Inferring pip requirements by reloading the logged model from the databricks artifact repository, which can be time-consuming. To speed up, explicitly specify the conda_env or pip_requirements when calling log_model().
```

```
print(accuracy_1)
```

```
# To get the best parameters
print("Tuned parameters for the best model: ")
print("MaxDepth: %g" %best_model_1.getMaxDepth())
print("MaxBins: %g" %best_model_1.getMaxBins())
print("Impurity: %s" %best_model_1.getImpurity())
```

```
# RandomForest Classifier
```

```
from pyspark.ml.classification import RandomForestClassifier
```

```
2023/05/29 13:12:35 INFO mlflow.utils.autologging_utils: Created MLflow autologging run with ID '35931319b9124ed586990c5576539974', which will track hyperparameters, performance metrics, model artifacts, and lineage information for the current pyspark.ml workflow
```

```
# Making predictions with the Random Forest classifier model
predictions2 = model2.transform(testingDF2)
predictions2.show()
```

```
|0.0265625|0.0296875|0.0278125| 0.0275| 0.02875| 0.035| 0.029375|0.0284375| 0.028125|0.0340625| 0.030625|0.0315625|
0.0265625|0.0253125|0.0259375|0.0284375|0.0296875|0.0340625|0.0328125|0.0328125| 1|[0.0265625,0.0296...| 1.0
|[19.0147178961624...|[0.95073589480812...| 0.0|
|0.0265625|0.0296875|0.0278125|0.0284375|0.0271875|0.0284375|0.0309375|0.0309375| 0.02875|0.0340625|0.0296875| 0.03125|
0.0359375| 0.671875|0.0640625| 0.07|0.0603125|0.0778125|0.0984375| 0.07875| 1|[0.0265625,0.0296...| 1.0
|[17.2870471161080...|[0.86020725588544...| 0.0|
```

```
#predictions2.display()
```

```
# Evaluating accuracy of the RandomForestClassifier without hyperparameter tuning
evaluator = MulticlassClassificationEvaluator(labelCol="label", predictionCol="prediction", metricName="accuracy")
accuracyRF = evaluator.evaluate(predictions2)
print("Accuracy score: ", accuracyRF)

Accuracy score: 0.9777979157227005
```

```
# Random Forest classifier model using several hyperparameters and getting the best model
with mlflow.start_run(run_name="random_forest_classification") as run:
```

```
    # create an instance of Random Forest classifier
    rf = RandomForestClassifier(labelCol="label", featuresCol="features")

    # parameter grid with different values
    RFParamGrid = (ParamGridBuilder()
                    .addGrid(rf.numTrees, [10,15 ])
                    .addGrid(rf.maxDepth, [9,13])
                    .build())

    # create a TrainValidationSplit object
    RFtvs = TrainValidationSplit(estimator=rf,
                                estimatorParamMaps=RFParamGrid,
                                evaluator=MulticlassClassificationEvaluator()
                                )

    # fit the model and retrieve the best model
    model = RFtvs.fit(trainingDF2)

    # retrieving a dictionary of hyperparameters and values of the best model
    best_parameters = model.bestModel.extractParamMap()
    # log the best model parameters
    for param, value in best_parameters.items():

        with mlflow.start_run(nested=True):
            mlflow.log_param("depth", 9)
        with mlflow.start_run(nested=True):
            mlflow.log_param("depth", 13)

    # The best trained model
    best_model_2 =model.bestModel

    # using the best model on the testing data

    evaluator = MulticlassClassificationEvaluator(labelCol="label",
                                                  predictionCol="prediction",
                                                  metricName="accuracy")
    accuracy_2 = evaluator.evaluate(best_model_2.transform(testingDF2))

    # log the accuracy of the best model
    mlflow.log_metric("accuracy", accuracy_2)

    # end the MLflow run
    mlflow.end_run()
```

2023/05/29 13:41:23 INFO mlflow.spark: Inferring pip requirements by reloading the logged model from the databricks artifact repository, which can be time-consuming. To speed up, explicitly specify the conda_env or pip_requirements when calling log_model().

```
2023/05/29 13:42:36 INFO mlflow.spark: Inferring pip requirements by reloading the logged model from the databricks artifact repository, which can be time-consuming. To speed up, explicitly specify the conda_env or pip_requirements when calling log_model().
```

```
# Accuracy using the best model
print(accuracy_2)
```

```
0.9841413683733575
```

```
# Parameters of best model from Random forest classifier
print("Tuned parameters for the best model: ")
print("MaxDepth: ", best_model_2.getMaxDepth())
print("numTrees: ", best_model_2.getNumTrees)
```

```
Tuned parameters for the best model:
```

```
MaxDepth: 9
```

```
numTrees: 15
```